

# Intel Image Classification

## End-to-End Deep Learning Pipeline

Ali Traore

12. Januar 2026

# Agenda

- 1 Einführung
- 2 Datensatz
- 3 Modellarchitekturen
- 4 Training & Evaluation
- 5 API & Deployment
- 6 Ergebnisse
- 7 Lessons Learned
- 8 Ausblick
- 9 Zusammenfassung

## Ziel

End-to-End Bildklassifikationssystem für 6 Landschaftskategorien mit Production-ready API und Dashboard

### Kategorien:

- Buildings
- Forest
- Glacier
- Mountain
- Sea
- Street

### Tech Stack:

- TensorFlow/Keras
- FastAPI + Uvicorn
- Streamlit + Plotly
- Docker Compose

## Vollständige ML Pipeline:

- ① **Datenverarbeitung:** Data loading, augmentation, preprocessing
- ② **Modelltraining:** 3 CNN-Architekturen mit verschiedenen Strategien
- ③ **Evaluation:** Metriken, Visualisierungen, Grad-CAM
- ④ **API Backend:** FastAPI mit automatischer Dokumentation
- ⑤ **Frontend:** Multi-page Streamlit Dashboard
- ⑥ **Deployment:** Docker Compose für einfache Bereitstellung

## Intel Image Classification:

- ~25.000 Bilder
- $150 \times 150$  RGB
- 6 Klassen (balanciert)

## Split:

- Train: ~14.000
- Validation: ~3.000
- Test: ~7.000

## Preprocessing:

- Normalisierung [0,1]
- Data Augmentation:
  - Rotation ( $\pm 20^\circ$ )
  - Horizontal Flip
  - Zoom (10%)
  - Shift (10%)
- Batch Size: 32

Modell	Parameter	Val Accuracy	Strategie
Baseline CNN	~2M	73%	Simple CNN
Regularized CNN	~2M	79%	+ L2 + Dropout
MobileNetV2	~3.5M	85-90%	Transfer Learning

## Implementierung:

- `src/models/`: `baseline.py`, `regularized.py`, `transfer_learning.py`
- `src/training/`: Unified trainer + callbacks
- `scripts/train_pipeline.py`: Komplette Pipeline

## Baseline CNN:

- 3x Conv2D + MaxPooling
- Flatten + Dense
- Dropout (0.5)
- Problem: Overfitting (73% val)

## Regularized CNN:

- + Batch Normalization
- + L2 Regularization
- + Dropout (0.6)
- Besser: 79% val

## MobileNetV2:

- Pre-trained (ImageNet)
- Custom Head:
  - GlobalAvgPooling
  - Dense(256)
  - Dropout(0.5)
  - Dense(6)
- Beste: 85-90% val

## Training Features:

- EarlyStopping (patience=10)
- ModelCheckpoint
- ReduceLROnPlateau
- TensorBoard Logging
- CSV Metrics Export

## Implementierung:

- src/training/trainer.py
- src/training/callbacks.py
- scripts/train\_pipeline.py

## Evaluation:

- Accuracy, Precision, Recall, F1
- Confusion Matrix
- Training/Validation Curves
- Grad-CAM Visualisierung

## Outputs:

- outputs/models/\*.keras
- outputs/figures/
- outputs/logs/

## RESTful API (api/main.py):

### Endpoints:

- GET / - Health Check
- GET /models - List Models
- POST /predict/{model} - Classify
- POST /feedback - User Feedback

### Features:

- Lazy Model Loading
- Model Caching
- CORS Support
- Auto Documentation (Swagger)

### Modelle:

- baseline
- regularized
- transfer\_learning

### Deployment:

- Uvicorn ASGI Server
- Port 8000
- Docker Container
- Health Checks

## Multi-Page Dashboard (streamlit/):

### Seiten:

- `_Home.py`: Upload & Predict
- `_Model_Comparison.py`: Vergleich
- `_Data_Analysis.py`: Dataset Stats
- `_Feedback_Dashboard.py`: Analytics

### Components:

- `status_panel.py`
- `visualization.py`

### Features:

- Drag & Drop Upload
- Model Selection
- Confidence Scores
- Grad-CAM Heatmaps
- Confusion Matrices
- Plotly Charts
- Feedback System

Port: 8501

# Docker Deployment

## Docker Compose Setup:

### Services:

- api: FastAPI (Port 8000)
- streamlit: Dashboard (Port 8501)

### Features:

- Health Checks
- Volume Mounts
- Auto Restart
- Service Dependencies

### Usage:

- docker compose up --build
- docker compose up -d
- docker compose logs -f
- docker compose down

### Vorteile:

- Einfaches Deployment
- Reproduzierbar
- Portabel

# Ergebnisse

Modell	Train Acc	Val Acc	F1-Score
Baseline	82%	73%	0.71
Regularized	84%	79%	0.78
MobileNetV2	91%	88%	0.88

## Beste Performance: MobileNetV2 Transfer Learning

- Validation Accuracy: 88.5%
- Alle Klassen: 85-92% Accuracy
- Herausforderung: Glacier vs. Mountain Verwechslung

## Technische Erkenntnisse:

- **Transfer Learning:** MobileNetV2 deutlich besser als Custom CNNs
- **Regularisierung:** Batch Norm + L2 + Dropout reduziert Overfitting
- **Data Augmentation:** Essentiell für Generalisierung
- **FastAPI:** Schnelle API-Entwicklung mit Auto-Docs
- **Streamlit:** Einfache Dashboard-Erstellung
- **Docker Compose:** Vereinfacht Multi-Service Deployment

## Mögliche Erweiterungen:

### Features:

- Ensemble Models
- Batch Prediction
- Video Classification
- Real-time Webcam

### Production:

- Cloud Deployment
- CI/CD Pipeline
- Monitoring (Prometheus)
- Model Versioning (MLflow)

## Projektergebnisse:

- 3 CNN-Modelle: Baseline, Regularized, Transfer Learning
- Beste Accuracy: 88.5% (MobileNetV2)
- FastAPI Backend mit 4 Endpoints
- Multi-Page Streamlit Dashboard
- Docker Compose Deployment
- Grad-CAM Interpretierbarkeit
- Komplette Dokumentation (README, RUN.md)

## Fazit

End-to-End ML Pipeline: Von Daten bis Production-Deployment

## Code Organization:

### Core ML:

- src/models/
- src/training/
- src/evaluation/
- src/utils/
- scripts/train\_pipeline.py

### Deployment:

- api/main.py
- streamlit/app.py
- streamlit/pages/
- Dockerfile
- docker-compose.yml

# Vielen Dank!

Fragen & Diskussion

Ali Traore

**Demo:** docker compose up --build

**API:** <http://localhost:8000/docs>

**Dashboard:** <http://localhost:8501>